

LAB 1: AI-ASSISTED CITATION INTEGRITY AUDIT

Generate a Research Paper with AI, Then Audit Its Citations

AI Tools for Research • 2 Hours • Individual Activity • Topic IDs T1-T40

Student Detail	Fill in
Name	Dharavathu Sravanthi
Roll Number	MNS2026011
Programme / Department	MTech. IT (Network system and computing)
Topic ID	23
Full Assigned Topic	The influence of citation formatting on human trust in fabricated resources
AI Tool Used	Chat GPT
Date	21-08-2026

IMPORTANT: You will receive a topic from T1-T40. Your task is not to produce the “best” paper. Your task is to generate an AI-written research paper and systematically determine whether its citations can be trusted.

1. PURPOSE OF THIS LAB

Generative AI systems can produce research papers that look scholarly and may contain convincing references, author names, journal titles, DOIs, arXiv IDs, and in-text citations. However, the presence of a citation does not guarantee that the cited publication actually exists, that its metadata are correct, that the identifier belongs to that publication, or that the cited paper supports the statement made by the AI.

In this laboratory activity, you will test these issues yourself by generating a research paper and auditing a systematic sample of its references.

Question 1 - Citation Authenticity: Does the cited publication actually exist, and is it correctly identified?

Question 2 - Claim-Citation Support: If the publication is genuine, does it actually support the claim for which the AI cited it?

2. IMPORTANT RULES

- Use only the topic assigned to you by the instructor.
- Generate the paper in a fresh AI conversation.
- Do not tell the AI that you are conducting a citation audit or hallucination test.
- Do not intentionally ask the AI to generate fake references.
- Preserve the original AI-generated paper exactly as produced.
- Do not correct, regenerate, replace, or question references before saving the original output.
- For verification, do not accept another AI chatbot's statement as proof that a reference exists.
- Use scholarly databases, publisher pages, DOI/arXiv/PMID records, or other reliable scholarly sources.
- Record what you actually find, even if all references are genuine or all are problematic.

Most important principle: You are not being evaluated on whether the AI produces good or bad references. You are being evaluated on how systematically and correctly you audit them.

3. WHAT YOU WILL DO

Part	Activity
A	Record your baseline perception of AI citations
B	Record the AI environment used

C	Generate and preserve an AI-written research paper
D	Create a reference inventory
E	Select 10 references systematically
F	Judge their plausibility before verification
G	Verify authenticity, metadata and identifiers
H	Determine whether genuine references support the AI-generated claims

4. RECOMMENDED 2-HOUR TIME PLAN

Time	Activity
0-10 min	Parts A-B: Baseline and AI environment
10-30 min	Part C: Generate and save research paper
30-40 min	Parts D-E: Inventory and select references
40-45 min	Part F: Pre-verification plausibility
45-80 min	Part G: Authenticity and metadata audit
80-110 min	Part H: Claim-citation entailment audit
110-120 min	Calculate scores, check worksheet and prepare submission

Do not spend excessive time formatting the AI-generated paper. The audit is the main activity.

PART A - BASELINE ASSESSMENT

Complete this section BEFORE opening the AI tool.

A1. How often do you use generative AI for academic work?

Never Rarely Occasionally Weekly ☒ Almost daily

A2. Have you previously asked an AI system to generate references?

Frequently ☒ Occasionally Once or twice Never

A3. Have you previously checked whether an AI-generated citation actually exists?

Always Sometimes ☒ Rarely Never Not applicable

A4. Before this lab, how confident are you that citations generated by AI are generally reliable?

1 2 3 ☒ 4 5 (1 = not at all confident; 5 = completely confident)

Important: Do not change this answer later, even if your opinion changes during the experiment.

PART B - RECORD YOUR AI ENVIRONMENT

Information	Your Response
AI tool used	Chat GPT
Model name	GPT-5.6 Luna
Model/version shown	Gpt-5.6-luna
Access tier	Free
Web browsing available?	Yes
Was web browsing actually used?	Yes
Research/Deep Research mode used?	Yes
Date/time generated	21-08-2026 10: 32 am

Important: If the model/version is not displayed, write NOT SHOWN. Do not guess the model version.

PART C - GENERATE THE AI RESEARCH PAPER

C1. Open a Fresh AI Conversation

Do not tell the AI that this is a citation audit, that you are testing hallucinations, or that you expect fabricated references. We want to observe normal AI research-paper generation behaviour.

C2. Use the Following Common Prompt

Replace only [YOUR ASSIGNED TOPIC].

Generate a scholarly research paper on the topic: "[YOUR ASSIGNED TOPIC]". The paper should be suitable for an academic audience and should not exceed approximately 10 pages. Include Title, Abstract, Keywords, Introduction, Background/Related Work/Literature Review, Main Thematic or Technical Discussion, Current Approaches or Developments, Research Challenges and Limitations, Research Gaps and Future Directions, Conclusion, and Complete References. Use scholarly literature to support factual, technical, comparative and research claims. Provide in-text citations throughout the paper. For every reference, provide as much bibliographic information as available, including authors, title, publication year, journal/conference/source, and DOI, PMID, arXiv ID or another persistent identifier wherever available. Use genuine scholarly sources. Do not intentionally fabricate references.

C3. Preserve the Original Output

Save the complete AI-generated paper

Save the complete reference list

Take screenshot(s) showing the generated output/references

Save conversation/export/share evidence if available

STOP: Do not ask the AI to correct or verify anything yet. Your original output is the experimental evidence.

PART D - PAPER PROFILE AND REFERENCE INVENTORY

Paper title: *The Influence of Citation Formatting on Human Trust in Fabricated Resources*

Measure	Value
Approximate pages	10
Approximate word count	4500 - 5000
Number of sections	11
Total number of references	11
Approximate in-text citation occurrences	35-40
Did AI warn you to verify references?	No

D1. Reference Inventory

Classify every reference once using the best available identifier in this priority order: DOI -> PMID -> arXiv ID -> URL only -> No persistent identifier. The category totals must equal the total number of references in the AI-generated paper.

Reference Type	Count
DOI available	17
PMID available, no DOI	0
arXiv ID available, no DOI/PMID	2
URL only	1
No persistent identifier supplied	0
TOTAL	20

Examples

Category	Example	How to Count
----------	---------	--------------

DOI available	Kather, J. N., et al. (2019). Predicting survival from colorectal cancer histology slides using deep learning. PLoS Medicine. DOI: 10.1371/journal.pmed.1002730	DOI available
PMID available, no DOI	A biomedical paper for which the AI provides PMID: 12345678, but no DOI is supplied/found	PMID available, no DOI
arXiv ID available	Yao, S., et al. ReAct: Synergizing Reasoning and Acting in Language Models. arXiv:2210.03629	arXiv ID available
URL only	World Health Organization. Colorectal cancer. Available at: https://... with no DOI, PMID or arXiv ID	URL only
No persistent identifier	Goodfellow, I., Bengio, Y., & Courville, A. (2016). Deep Learning. MIT Press.	No persistent identifier if none is supplied

Important: Absence of a DOI does not mean a reference is fake. An identifier being present also does not prove authenticity; it must be checked.

PART E - SELECT 10 REFERENCES FOR AUDIT

You are not allowed to choose only suspicious-looking references. This prevents selection bias.

If the bibliography contains 10 or fewer references: audit all references.

If it contains more than 10 references: select the first 3 references, the last 3 references, and 4 references approximately evenly distributed through the middle.

Example: If there are 28 references, a reasonable sample could be 1, 2, 3, 8, 13, 18, 23, 26, 27, 28.

Audit Slot	Original Bibliography Reference No.
R1	1
R2	2
R3	3
R4	17
R5	18
R6	19
R7	7
R8	10
R9	13
R10	15

Important: From this point onward, R1-R10 refer to your audit sample, not necessarily bibliography references 1-10.

PART F - PLAUSIBILITY TEST

Do this BEFORE searching for any reference.

Look only at the citation as generated by the AI. Ask yourself: “If I had not verified this citation, how convincing would it look to me?”

Score	Meaning
1	Obviously suspicious
2	Somewhat suspicious
3	Unsure
4	Looks genuine
5	Completely convincing

Ref.	Initial Impression	Plausibility 1-5	Predict Genuine?
R1	Looks convincing	5	Yes
R2	Completely convincing	5	Yes
R3	Looks genuine	5	Yes

R4	Looks genuine	4	Yes
R5	Looks convincing	5	Yes
R6	Looks genuine	4	Yes
R7	Looks convincing	5	Yes
R8	Looks convincing	5	Yes
R9	Looks genuine	4	Yes
R10	Looks genuine	4	Yes

Example: A citation with realistic authors, a well-known journal, complete volume/pages and a DOI-like identifier may look highly convincing. You might record Plausibility = 5 and Predict Genuine = Yes. You are recording your pre-verification belief, not the truth.

Critical rule: Never change Part F after completing Part G. Your incorrect predictions are valuable experimental observations.

PART G - CITATION AUTHENTICITY AND METADATA AUDIT

G1. What Are You Checking?

- Does the publication exist?
- Is the title correct?
- Are the authors correct?
- Is the year correct?
- Is the journal/conference/source correct?
- Are volume, issue and pages correct, if supplied?
- Does the DOI/arXiv/PMID belong to this publication?

G2. Verification Procedure

Persistent identifier: DOI / arXiv / PMID

Publisher or official repository

Google Scholar

Semantic Scholar / OpenAlex

Crossref

PubMed or another discipline-specific database

Exact-title search

Author + title-keyword search

Not acceptable as sole evidence: “I asked another AI and it said the paper exists.” AI can help you search, but final verification must rely on scholarly or publisher evidence.

G3. Assign One A-E Code

Code	Classification	Meaning
A	VERIFIED	Publication exists and important metadata are correct.
B	WRONG METADATA	Publication exists, but one or more bibliographic details are incorrect.
C	FRANKENSTEIN	Real-looking pieces are combined, but that specific publication does not exist.
D	FABRICATED	No reliable scholarly evidence that the publication exists.
E	IDENTIFIER MISMATCH	DOI/arXiv/PMID is invalid or resolves to a different publication.

G4. Examples

A: AI gives a paper and all details match the publisher record.

B: The paper exists, but the AI gives the wrong publication year or journal.

C: The authors are real, the journal is real, and the topic is plausible, but no such paper with that title/authorship combination exists.

D: You cannot find the paper or a credible scholarly record after systematic searching.

E: The DOI supplied by AI opens an entirely different paper or does not resolve.

G5. Record Your Evidence

Ref.	Publication Found?	Title Match?	Authors Match?	Year Match?	Venue Match?	Identifier Match?
R1	Y	Y	Y	Y	Y	Y
R2	Y	Y	Y	Y	Y	Y
R3	Y	Y	Y	Y	Y	Y
R4	N	N	N	N	N	NA
R5	Y	Y	Y	Y	Y	Y
R6	Y	Y	Y	Y	Y	Y
R7	Y	Y	Y	Y	Y	Y
R8	Y	Y	Y	Y	Y	Y
R9	Y	Y	Y	Y	Y	Y
R10	Y	Y	Y	Y	Y	Y

Ref.	Verification Source(s)	Identifier Checked	What You Found
R1	Publisher page	Yes	Publication exists and important metadata are correct.
R2	Publisher page	Yes	Publication exists and important metadata are correct.
R3	Publisher page	yes	Publication exists and important metadata are correct.
R4	FABRICATED	yes	No reliable scholarly evidence that the publication exists
R5	Article	Yes	Publication exists and important metadata are correct.
R6	Article	yes	Publication exists and important metadata are correct.
R7	Public record	Yes	Publication exists and important metadata are correct.
R8	Public record	Yes	Publication exists and important metadata are correct.
R9	Paper	Yes	Publication exists and important metadata are correct.
R10	Paper	yes	Publication exists and important metadata are correct.

G6. Final Classification

Ref.	Final Code	One-Line Reason
R1	A/B/C/D/E	A
R2	A/B/C/D/E	A
R3	A/B/C/D/E	A
R4	A/B/C/D/E	D
R5	A/B/C/D/E	A
R6	A/B/C/D/E	A
R7	A/B/C/D/E	A
R8	A/B/C/D/E	A
R9	A/B/C/D/E	A
R10	A/B/C/D/E	A

Good reason: B: Article exists, but AI reported 2022 instead of 2020.

Poor reason: B: Details wrong.

PART G - SCORING

Count your final classifications: A = 9____ B = ____0 C = 0____ D = 1____ E = 0____

Total audited: $N = A + B + C + D + E = 10$

Code	Points
A	4
B	3
C	1
D	0
E	1

Authenticity Points = $4A + 3B + C + E = 36$

Authenticity Score = $[(4A + 3B + C + E) / (4N)] \times 100 = 90$

My calculation: $4(9) + 3(0) + 1(0) + 0(\underline{0}) + 1(0) = 36$ points Authenticity Score

= $\frac{36}{(4 \times 10)} \times 100 = 90$

Worked example: If A=6, B=2, C=1, D=1, E=0, then points = $4(6)+3(2)+1(1)=31$; maximum = 40; Authenticity Score = $31/40 \times 100 = 77.5$.

Score	Interpretation
90-100	Excellent authenticity
75-89	Good, minor concerns
60-74	Moderate citation-authenticity risk
40-59	High citation-authenticity risk
Below 40	Very high citation-authenticity risk

G7. Prediction Accuracy

Compare your Part F Genuine Yes/No predictions with the Part G results.

Prediction Accuracy = $(\text{Correct genuine/fake predictions} / \text{Total predictions}) \times 100$

Correct predictions = 9 Total predictions = 10 Prediction Accuracy = 90%

FINAL LAB 1 RESULTS

Measure	Your Result
Total references in AI paper	10
References deeply audited	10
A: Verified	9
B: Wrong metadata	0
C: Frankenstein	0
D: Fabricated	1
E: Identifier mismatch	0
Authenticity Score	90/100
Prediction Accuracy	90%

FINAL INTERPRETATION

1. Did professional-looking citations necessarily turn out to be genuine?

.....In my view its not necessary to look like that, some may look professional but fabricated.....

.....

2. Did all genuine references actually support the claims for which they were cited?

.....

.....Most of the audited information claims to be genuine from AI.....

3. What was the most serious citation-integrity problem you observed?

.....

As for the time being I haven't seen any serious citation problem. With my audit being done I found 90 percent to be correct info retrieved.....

4. What is the single most important lesson you learned from this lab?

.....

.....I have developed any subconscious thing of AI can even do good work with research things not only embodying to the academics and computation.....

SUBMISSION INSTRUCTIONS

Submit one PDF: Lab1_RollNumber_Name

Example: Lab1_IEC2026001_RahulSharma

STUDENT DECLARATION

I confirm that I generated the research paper using the recorded AI environment, preserved the original output before conducting the audit, followed the prescribed reference-selection procedure, and personally verified the evidence reported in this worksheet. I have not altered my pre-verification predictions after seeing the verification results.

Student Name:Dharavathu Sravanthi.....

Roll Number:MNS2026011.....

Signature:Sravanthi.....

Date:21-08-26.....

CENTRAL LESSON OF LAB 1

Does the source exist? + Does the source actually support the claim?

AI Generation -> Reference Inventory -> Systematic Sampling -> Prediction -> Authenticity Verification -> Claim-Support
Verification -> Quantitative Scores -> Submission

INTERACTION WITH AI (CHATGPT):

<https://chatgpt.com/share/6a88661f-1ce0-83e8-9465-70a1e6a767c1>

REFERENCES TAKEN:

Yes. I checked the bibliography against scholarly/publisher records and expanded the literature base with **newer 2025–2026 work**. I would recommend using the expanded set rather than simply adding references indiscriminately: several of the newer papers are especially relevant to the paper's central argument about **fabricated citations, surface plausibility, verification, and trust**.

One important correction from the earlier paper: some of the newest studies were not available when the original draft was written, so they should be incorporated into a revised literature review rather than merely appended to the bibliography.

1. Original resources used in the paper

1. Walters & Wilder (2023) — Essential

Walters, W. H., & Wilder, E. I. (2023).

Fabrication and errors in the bibliographic citations generated by ChatGPT.

Scientific Reports, 13, 14045.

DOI: **10.1038/s41598-023-41032-5**

PMID: **37679503**

PMCID: **PMC10484980**

This is probably the **single most important paper** for your topic. It directly investigates fabricated bibliographic references generated by ChatGPT. The study examined 636 citations from 84 AI-generated literature reviews and found fabrication rates of 55% for GPT-3.5 and 18% for GPT-4. ([DOI](#))

[Read the full Scientific Reports paper](#)

[Read the PubMed record](#)

2. Ding et al. (2025) — Essential for the trust component

Ding, Y., Facciani, M., Joyce, E., Poudel, A., Bhattacharya, S., Veeramani, B., Aguinaga, S., & Weninger, T. (2025).

Citations and Trust in LLM Generated Responses.

Proceedings of the AAAI Conference on Artificial Intelligence, 39(22), 23787–23795.

DOI: **10.1609/aaai.v39i22.34550**

arXiv: **2501.01303**

This is the strongest source connecting **citations directly to human trust in AI-generated information**. Participants trusted responses more when citations were present, even when citations were randomly selected; checking citations reduced trust. ([AAAI Publications](#))

[AAAI paper page](#)

[arXiv version](#)

3. Hovland & Weiss (1951)

Hovland, C. I., & Weiss, W. (1951).

The Influence of Source Credibility on Communication Effectiveness.

Public Opinion Quarterly, 15(4), 635–650.

DOI: 10.1086/266350

This is a foundational source for the **source-credibility theory** underlying the argument that people use credibility cues when evaluating information. ([DOI](#))

[Oxford Academic record](#)

4. Fogg & Tseng (1999)

Fogg, B. J., & Tseng, H. (1999).

The Elements of Computer Credibility.

In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems (CHI '99)*, 80–87. ACM.

DOI: 10.1145/302979.303001

This provides the theoretical bridge from traditional credibility research to **computer-mediated information and interface cues**. ([Stanford Web Credibility Project](#))

[ACM DOI record](#)

5. Lucassen & Schraagen (2013)

Lucassen, T., & Schraagen, J. M. (2013).

The influence of source cues and topic familiarity on credibility evaluation.

Computers in Human Behavior, 29(4), 1387–1392.

DOI: 10.1016/j.chb.2013.01.036

Very useful for your discussion of **reader expertise, familiarity, and credibility cues**. The experiment found that source cues mattered particularly when the credibility of the information itself was ambiguous. ([ScienceDirect](#))

[ScienceDirect article](#)

6. Kang, Bae, Zhang & Sundar (2011)

Kang, H., Bae, K., Zhang, S., & Sundar, S. S. (2011).

Source Cues in Online News: Is the Proximate Source More Powerful than Distal Sources?

Journalism & Mass Communication Quarterly, 88(4), 719–736.

DOI: 10.1177/107769901108800403

Useful for your argument that **user involvement changes how source cues affect credibility**. ([Sage Journals](#))

[SAGE article](#)

7. Agarwal et al. (2023)

Agarwal, A., Arafa, M., Avidor-Reiss, T., Hamoda, T. A.-A. A.-M., & Shah, R. (2023).
Citation Errors in Scientific Research and Publications: Causes, Consequences, and Remedies.
The World Journal of Men's Health, 41(3), 461–465.
 DOI: [10.5534/wjmh.230001](https://doi.org/10.5534/wjmh.230001)
 PMID: [37118953](https://pubmed.ncbi.nlm.nih.gov/37118953/)
 PMCID: [PMC10307651](https://pubmed.ncbi.nlm.nih.gov/PMC10307651/)

Useful for establishing that citation problems are not exclusively an AI phenomenon and that citation accuracy has consequences for scholarly communication.

[PubMed record](#)

8. Besançon et al. (2024)

Besançon, L., Cabanac, G., Labbé, C., & Magazinov, A. (2024).
Sneaked references: Fabricated reference metadata distort citation counts.
Journal of the Association for Information Science and Technology, 75(12), 1368–1379.
 DOI: [10.1002/asi.24896](https://doi.org/10.1002/asi.24896)
 arXiv: [2310.02192](https://arxiv.org/abs/2310.02192)

Extremely useful because it expands your topic beyond AI-generated text to **manipulation of scholarly citation metadata itself**.

[Wiley article](#)

[arXiv version](#)

9. Camp, Bengtson & Sandstrom (2025)

Camp, N. T., Bengtson, J. A., & Sandstrom, J. C. (2025).
The citation catastrophe: Propagation of AI-generated counterfeit citations in scholarship.
The Journal of Academic Librarianship, 51(4), 103065.
 DOI: [10.1016/j.acalib.2025.103065](https://doi.org/10.1016/j.acalib.2025.103065)

This is especially valuable for your **future-directions section**, because it addresses how fabricated citations can propagate through subsequent scholarship.

10. Topaz et al. (2026) — Very important new evidence

Topaz, M., Roguin, N., Gupta, P., Zhang, Z., & Peltonen, L.-M. (2026).
Fabricated citations: An audit across 2.5 million biomedical papers.
The Lancet, 407(10541), 1779–1781.
 DOI: [10.1016/S0140-6736\(26\)00603-3](https://doi.org/10.1016/S0140-6736(26)00603-3)
 PMID: [42107362](https://pubmed.ncbi.nlm.nih.gov/42107362/)

This is one of the most important additions for a current version of your paper. It moves the discussion from "LLMs sometimes make up references" to **fabricated references appearing at scale in the scholarly literature**.

The audit examined approximately 97.1 million references and identified 4,046 fabricated references across 2,810 papers.

[\(MDPI\)](#)

[PubMed record](#)

11. Resnik & Hosseini (2026)

Resnik, D. B., & Hosseini, M. (2026).

Hallucinated citations produced by generative artificial intelligence may constitute research misconduct when citations function as data in scholarly papers.

Accountability in Research.

DOI: **10.1080/08989621.2026.2645390**

PMID: **41833014**

This is particularly useful for your **research-integrity and ethical implications** section. ([Taylor & Francis Online](#))

[Taylor & Francis article](#)

2. Additional papers I strongly recommend adding

I found several **newer 2025–2026 papers** that can substantially strengthen your research paper.

12. Davis & Mahmoud (2026) — Citation constraints and hallucinations

Davis, K., & Mahmoud, Q. (2026).

Citation Constraints and Reference Hallucinations in Large Language Models.

Proceedings of the 39th Canadian Conference on Artificial Intelligence, Proceedings of Machine Learning Research, 318, 464–475.

This study is directly relevant to your topic. It tested **Gemini 3, ChatGPT 5.1, ChatGPT 4o, and Microsoft 365 Copilot**, generating academic-style documents and verifying citations against Crossref, OpenAlex, and arXiv. It found that stricter citation requirements could actually be associated with higher rates of invalid or inconsistent references. ([Proceedings of Machine Learning Research](#))

[PMLR full paper](#)

Why I recommend it: It gives your paper a **2026 empirical source** connecting prompting/citation requirements with reference hallucination.

13. Misra & Udandaraao (2026) — Citation hallucination detection

Misra, N., & Udandaraao, V. (2026).

Detecting Citation Hallucinations in Large Language Model Outputs.

Proceedings of the AAAI Conference on Artificial Intelligence, 40(48), 41325–41327.

DOI: **10.1609/aaai.v40i48.42257**

The authors propose a hybrid citation-hallucination detection pipeline involving bibliographic retrieval, fuzzy similarity, and LLM-based verification. ([AAAI Publications](#))

[AAAI paper page](#)

Why I recommend it: Excellent for your **Current Approaches / Detection** section.

14. "Fluency Without Fidelity" (2026) — Highly relevant to your topic

Fluency Without Fidelity: Errors in Citation-Attributed Claims in Large Language Model-Generated Literature Reviews in Mental Health. (2026).

Journal of Technology in Behavioral Science.

The study investigated whether claims in LLM-generated literature reviews actually correspond to the scientific evidence in the cited sources. It used ChatGPT-5 to generate literature reviews and examined citation–claim alignment. ([Springer Link](#))

[Springer article](#)

Why I recommend it: Your paper currently emphasizes whether a citation *exists*. This paper lets you address the deeper issue of whether a **real citation actually supports the claim**.

15. Zhao et al. (2026) — Large-scale "in the wild" evidence

Zhao, Z., Wang, Y., Stuart, T., De Vaan, M., Ginsparg, P., & Yin, Y. (2026).

LLM hallucinations in the wild: Large-scale evidence from non-existent citations.

arXiv preprint, **arXiv:2605.07723**.

This is a particularly interesting new preprint. The authors audit **111 million references across 2.5 million papers** from arXiv, bioRxiv, SSRN, and PubMed Central and report a sharp increase in nonexistent references after widespread LLM adoption, with an estimate of approximately **146,932 hallucinated citations in 2025 alone**. ([arXiv](#))

[arXiv paper](#)

Important: Because this is an **arXiv preprint**, I would use it as emerging evidence rather than giving it the same evidentiary status as a peer-reviewed journal article.

16. Russinovich, Kumar & Salem (2026) — "Phantom References"

Russinovich, M., Siva Kumar, R. S., & Salem, A. (2026).

Phantom References: Hallucinated Citations That Survive Peer Review at Top-Tier Conferences.

arXiv preprint.

The study examines accepted papers from **ICLR, ICML, NeurIPS, and USENIX Security** and develops a citation-verification pipeline called **RefChecker**. It reports that hallucinated citations can survive peer review and describes a practical automated verification approach. ([arXiv](#))

[arXiv paper](#)

Why I recommend it: This is excellent for your argument that **peer review itself does not necessarily guarantee citation authenticity**.

17. Han et al. (2026) — Phantom citations

Han et al. (2026).

Phantom citations: An empirical study of non-existent and unverifiable references in scholarly literature.

Journal of the Association for Information Science and Technology.

DOI: [10.1002/asi.70104](https://doi.org/10.1002/asi.70104)

This is another directly relevant recent study concerning non-existent and unverifiable scholarly references and the surface plausibility of hallucinated citations. ([Wiley Online Library](#))

[Wiley article](#)

18. Alansari & Luqman (2026) — LLM hallucination review

Alansari, A., & Luqman, H. (2026).

Large language models hallucination: A comprehensive survey.

Computer Science Review.

DOI: [10.1016/j.cosrev.2026.100970](https://doi.org/10.1016/j.cosrev.2026.100970)

This is broader than citation hallucination specifically, covering causes, detection, and mitigation of LLM hallucinations. ([ScienceDirect](#))

[ScienceDirect article](#)

Why I recommend it: Useful for the theoretical background on why LLMs produce plausible but unsupported information.

19. 2026 citation-quality scoring study

A Large Language Model Pipeline for Automated Citation Quality Scoring Across Engineering Journal Quartiles with Expert Validation. (2026).

Scientific Reports.

The study analyzes approximately **2.2 million citations from 56,381 papers** and reports evidence of invalid/fabricated citations in human-authored engineering literature, while examining automated citation-quality scoring. It also reports a substantial verification gap among researchers and reviewers. ([Nature](#))

[Scientific Reports article](#)

20. Topaz et al. (2026) | Highly relevant to AI citation fabrication

Topaz, M., Roguin, N., Gupta, P., Zhang, Z., & Peltonen, L.-M. (2026). Fabricated citations: An audit across 2.5 million biomedical papers. *The Lancet*, 407(10541), 1779–1781. [https://doi.org/10.1016/S0140-6736\(26\)00603-3](https://doi.org/10.1016/S0140-6736(26)00603-3)
