

The Influence of Citation Formatting on Human Trust in Fabricated Resources

Abstract

Citations perform two related functions in scholarly communication: they identify the evidentiary basis of claims and provide readers with a route for verification. These functions make citation presentation potentially important to human trust, but they also create an opportunity for deception when bibliographic information is fabricated. This paper reviews interdisciplinary evidence on source credibility, information-processing heuristics, citation errors, bibliographic fabrication, and trust in citation-bearing artificial-intelligence (AI) outputs to examine how citation formatting may influence perceived trustworthiness even when the cited resource does not exist. The central argument is that formatting should be understood not merely as a technical convention but as a bundle of credibility cues. Author names, publication dates, journal titles, volume and issue information, page ranges, persistent identifiers, numbered reference systems, and polished typographic consistency can collectively signal scholarly legitimacy. Existing experimental research demonstrates that the mere presence of citations can increase trust in AI-generated answers, including under conditions in which citations are irrelevant, while research on source cues shows that credibility judgments are sensitive to contextual and heuristic signals. At the same time, direct experimental evidence isolating citation *formatting* from citation *presence*, *relevance*, and *source identity* remains limited. Recent audits demonstrate that fabricated references can be highly plausible and correctly formatted, and that fabricated references are increasingly appearing in the scholarly record. The paper therefore identifies a significant empirical gap: the need for controlled experiments that manipulate citation formatting independently of source validity. It proposes a research agenda combining factorial experiments, behavioral verification measures, expertise manipulations, eye tracking, and automated bibliographic verification. The paper concludes that polished formatting can plausibly increase initial trust in fabricated resources, but formatting alone should not be treated as a demonstrated causal determinant until directly tested.

Keywords: citation formatting; fabricated references; citation hallucination; source credibility; trust; information heuristics; scholarly communication; generative AI; bibliographic verification; misinformation

1. Introduction

Citations are fundamental infrastructure for scholarly communication. They allow readers to identify the intellectual provenance of claims, distinguish new contributions from prior work, and inspect the evidence underlying an argument. In principle, a citation therefore represents more than a typographic marker: it is an assertion that a particular source exists and that the source bears a meaningful relationship to the claim being made. When that source is nonexistent,

incorrectly identified, or irrelevant, the citation ceases to perform its evidentiary function while potentially retaining its persuasive appearance.

The problem has become particularly salient with the growth of generative artificial intelligence (GenAI). Large language models can generate bibliographies containing plausible authors, article titles, journals, publication years, page ranges, and digital object identifiers (DOIs), even when the purported works do not exist. Walters and Wilder (2023), for example, examined 636 references generated by GPT-3.5 and GPT-4 across 84 literature reviews and found substantial rates of fabricated citations, alongside substantial bibliographic errors among citations that corresponded to real works.

The credibility implications are consequential because readers rarely verify every citation. Citation evaluation competes with time constraints, information overload, and uneven subject expertise. Research on online credibility has repeatedly demonstrated that people use source and interface cues when evaluating information, particularly when the underlying information is difficult to assess directly (Fogg & Tseng, 1999; Lucassen & Schraagen, 2013).

Recent experimental work provides particularly relevant evidence. Ding et al. (2025) conducted an experiment involving AI-generated answers with zero, one, or five citations and found that citations increased perceived trustworthiness. Importantly, their experiment also manipulated citation validity: participants were exposed to both relevant and randomly selected citations. The study therefore demonstrates that citation presence itself can function as a trust cue, while citation checking can reveal skepticism.

The question addressed in this paper is narrower: **How does citation formatting influence human trust when the underlying resource is fabricated?** The distinction is important. Citation formatting includes conventions such as author–date presentation, numbered references, journal and volume information, DOI presentation, punctuation, italics, ordering, and bibliographic completeness. These features may contribute to a general impression of scholarly professionalism. However, a correctly formatted citation can still refer to a nonexistent article.

This paper argues that existing literature supports a theoretically plausible relationship between citation formatting and initial trust, but does not yet justify a strong causal claim that particular formatting styles independently make fabricated resources more believable. Instead, the evidence indicates that formatting is likely one component of a broader *credibility-cue system*. Establishing its independent effect requires experiments that separate formatting from source identity, citation relevance, textual quality, and actual verifiability.

2. Background and Related Work

2.1 Citation as an epistemic and social signal

The scholarly citation system serves several purposes simultaneously. Citations attribute intellectual credit, establish relationships among claims and prior scholarship, allow readers to locate supporting evidence, and contribute to the cumulative structure of academic knowledge.

Consequently, citation accuracy is not merely an editorial concern. Errors can alter the meaning of claims and propagate misinformation through subsequent publications (Agarwal et al., 2023).

Citation practices also contribute to institutional signals of legitimacy. A manuscript containing a substantial and apparently coherent bibliography can look more scholarly than an otherwise identical manuscript without references. This creates an asymmetry between *epistemic validity* and *epistemic appearance*: the former depends on whether sources actually support claims, whereas the latter can be influenced by visible features of presentation.

The distinction is especially important for fabricated references. A fabricated reference can be characterized as a bibliographic entry purporting to identify a scholarly work that does not exist. This differs from a citation error involving a real publication but incorrect page numbers, authorship, title, DOI, or interpretation. Walters and Wilder (2023) explicitly distinguish fabricated works from substantive errors in citations to real works, demonstrating that both problems can coexist in generated bibliographies.

2.2 Source credibility and heuristic processing

Research on persuasion has long established that perceived source characteristics can affect judgments of communication. Hovland and Weiss (1951) demonstrated experimentally that source credibility influences communication effectiveness, establishing an influential foundation for subsequent credibility research. Their work conceptualized credibility principally through perceptions of source expertise and trustworthiness.

Fogg and Tseng (1999) extended credibility analysis into computing environments, defining computer credibility around users' perceptions of believability and identifying mechanisms through which technological systems can acquire credibility. This work is relevant because citation presentation is encountered within a technological interface and can operate as a visible cue before a user independently verifies a source.

Later studies demonstrate that source cues can be influential under conditions of ambiguity. Lucassen and Schraagen (2013), for example, found that source cues affected credibility judgments particularly when the credibility of information itself was ambiguous. Their experiment using Wikipedia materials also showed that topic familiarity moderated how source cues influenced trust.

Similarly, Kang et al. (2011) found that source cues in online news affected credibility judgments differently depending on users' involvement: highly involved readers considered multiple source cues, whereas low-involvement readers were more strongly influenced by the proximate source. These findings suggest that citation-format effects should not be expected to operate uniformly. Users with substantial domain expertise or high motivation to verify evidence may be less susceptible to superficial formatting cues.

2.3 Citation presence and trust in AI-generated information

The emergence of AI-generated text provides an unusually direct setting for studying citation-based trust. Ding et al. (2025) experimentally compared AI-generated responses with different numbers and qualities of citations. Their findings indicate that citation presence increased perceived trustworthiness, while invalid or random citations reduced trust when users actually inspected them. The experiment also found no meaningful trust advantage for five citations over one citation, suggesting that citation quantity may exhibit diminishing returns.

This finding has direct implications for fabricated resources. If a reader initially treats a citation as a verification opportunity or social-proof cue, a fabricated citation may benefit from its *appearance* before its validity is checked. Once verification occurs, however, the credibility mechanism can reverse: a nonexistent source provides evidence against the credibility of the surrounding content.

This creates a two-stage model of citation-based trust:

1. **Pre-verification stage:** visible citation cues contribute to an initial credibility judgment.
2. **Post-verification stage:** source existence, relevance, and evidentiary correspondence determine whether trust is maintained, increased, or reduced.

The distinction is central to understanding why formatting could matter without being epistemically meaningful.

3. Citation Formatting as a Credibility Cue

3.1 What constitutes citation formatting?

Citation formatting encompasses the visible structure used to represent a source. Depending on the disciplinary system, this may include:

- author–date citations such as *(Smith, 2024)*;
- numbered citations such as *[12]*;
- footnotes or endnotes;
- full author names or abbreviated author lists;
- article and journal titles;
- volume, issue, and page information;
- publication dates;
- DOI or other persistent identifiers;
- alphabetical or numerical ordering;
- italics, capitalization, punctuation, and indentation;
- consistency across entries.

Different disciplines employ different conventions. Vancouver-style references generally use numerical citation systems, whereas Harvard-style systems use author–date conventions; individual journals then impose further variations in ordering and punctuation (Klein & Wager, 2019).

Importantly, there is little reason to assume that one legitimate scholarly style is intrinsically more truthful than another. Formatting systems are conventions rather than truth-detection mechanisms. The relevant psychological question is instead whether some formats *signal* greater legitimacy, professionalism, specificity, or verifiability.

3.2 Completeness and specificity

A citation containing author, year, journal, volume, issue, pages, and DOI appears more information-rich than a citation containing only an author and title. Such specificity may increase perceived legitimacy because it creates multiple apparent points of verification.

The paradox is that fabricated citations can exploit precisely this property. Topaz et al. (2026), in an audit of 2.5 million biomedical papers, found 4,046 fabricated references across 2,810 papers among approximately 97.1 million references that could be verified in their database-based analysis. The fabricated entries were not necessarily visually defective; the researchers reported that examples could be topically specific, correctly formatted, attributed to real researchers, and associated with plausible publication years.

This finding is critical because it weakens the assumption that formatting irregularity is a reliable indicator of fabrication. A fabricated citation may imitate the statistical and stylistic regularities of genuine scholarly references.

3.3 Consistency as a signal

Consistency may constitute another credibility cue. A bibliography in which every reference follows the same punctuation, capitalization, ordering, and DOI convention can convey editorial control. Conversely, inconsistent formatting can make a document appear hastily produced or unreliable.

However, this cue can be reproduced mechanically. A language model or automated reference generator can produce a bibliography with extremely consistent formatting without possessing knowledge of whether the underlying works exist. Consequently, formatting consistency can become detached from epistemic reliability.

This resembles a broader problem in digital credibility research: visual or interface characteristics can provide useful shortcuts when they correlate with quality, but they become exploitable when those correlations can be reproduced independently of the underlying property. Fogg and Tseng's (1999) credibility framework is useful here because it emphasizes the distinction between perceived credibility and actual validity.

3.4 Persistent identifiers and the appearance of verifiability

DOIs and other persistent identifiers have a particularly interesting status. A DOI normally functions as a verification mechanism: a reader can follow it to an identified scholarly object. Yet the presence of a DOI-like string can itself serve as a visual marker of scholarly legitimacy.

A fabricated DOI therefore illustrates the distinction between *verification affordance* and *verification outcome*. The identifier may appear highly authoritative before it is resolved. If it resolves to nothing, the credibility benefit disappears. If it resolves to another work, the discrepancy can be even more damaging because it reveals an identity-level mismatch.

The problem is not restricted to GenAI. Besançon et al. (2024) demonstrated that bibliographic metadata can be manipulated independently of the visible version of record. In their case study, "sneaked" references were inserted into DOI metadata and subsequently propagated to bibliometric platforms even though the references did not appear in the original articles. They identified at least 5,978 such references in their three-journal case corpus, representing approximately 9% of registered references in that subset.

This demonstrates that trust in citation infrastructure can be exploited at several levels: the textual bibliography, persistent identifiers, publisher metadata, and bibliometric databases.

4. Why Fabricated Citations Can Appear Trustworthy

The credibility of a fabricated resource can be understood as the product of several interacting signals rather than formatting alone.

4.1 Authority cues

Real-looking author names, prestigious institutional affiliations, established journal titles, and plausible publication dates can provide authority cues. Source-cue research indicates that perceived expertise and institutional identity can influence credibility judgments (Hovland & Weiss, 1951; Go et al., 2014). Go et al. (2014), for example, found that expertise cues in online news affected perceived credibility and interacted with other interface cues.

A fabricated citation that combines a recognizable journal with a plausible author name may therefore benefit from several mutually reinforcing cues.

4.2 Cognitive fluency

A second mechanism is processing fluency. Information that is easy to process may feel more familiar or plausible, although fluency is not equivalent to truth. In the citation context, consistent formatting reduces the cognitive effort required to interpret bibliographic information. A reader accustomed to APA or Vancouver conventions may process a correctly formatted entry automatically.

The resulting familiarity can create an important vulnerability: a fabricated citation that conforms to a reader's expectations may attract less scrutiny than an obviously malformed one.

4.3 Verification cost

Verification imposes a cost. Readers must click a DOI, search a database, inspect the publication record, and determine whether the cited work actually supports the claim. This takes time and requires sufficient bibliographic literacy.

Agarwal et al. (2023) argue that citation verification is frequently neglected and report numerous categories of citation error, including incorrect technical details, wrong sources, factual errors, and unjustified interpretations. Thus, the effectiveness of fabricated citations may depend partly on the probability that a reader will not perform independent verification.

4.4 Social proof

Citations also communicate that other researchers have supposedly produced relevant knowledge. Ding et al. (2025) explicitly interpret citations in AI-generated answers partly through a social-proof framework. Their experimental results demonstrate that citation presence can increase perceived trustworthiness even before the citation is checked.

A fabricated citation can therefore mimic two signals simultaneously: *someone authoritative supposedly published this* and *the current author has apparently documented the claim*. Formatting can reinforce both signals.

5. Current Approaches and Developments

5.1 Manual verification

The most direct defense against fabricated references remains manual verification. A reader can search bibliographic databases, resolve DOIs, inspect publisher records, and compare the cited claim with the source itself. This remains the strongest semantic safeguard because it can assess not only whether a source exists but whether it actually supports the claim.

Its weakness is scalability. Modern articles can contain hundreds of references, and verification requires substantial labor.

5.2 Reference-management and metadata tools

Reference managers and structured metadata services reduce conventional bibliographic errors. DOI lookup, PubMed identifiers, Crossref metadata, and citation-management software can help confirm author names, titles, dates, and publication venues.

However, these systems cannot automatically guarantee that the cited work supports the claim. Moreover, metadata themselves can be manipulated, as demonstrated by Besançon et al. (2024).

5.3 Automated citation verification

Recent research is moving toward automated reference-integrity systems. Walters and Wilder (2023) demonstrated systematic database checking of AI-generated bibliographies, while Topaz et al. (2026) developed a large-scale verification pipeline that compared reference metadata against multiple scholarly databases.

The scale of the latter study is particularly significant: its audit identified thousands of fabricated references embedded in peer-reviewed biomedical literature. The incidence increased substantially during the period examined, although the authors caution that the methodology cannot establish the precise causal contribution of AI versus other mechanisms such as paper mills or changes in publishing practices.

5.4 AI-assisted verification

AI itself can be used defensively. Automated systems can classify bibliographic anomalies, reconcile title variants, identify DOI mismatches, and prioritize suspicious references for human inspection.

This creates a recursive problem: AI-generated fabrication may require AI-assisted verification, but verification systems must themselves be evaluated for false positives and false negatives. The appropriate architecture is therefore not "AI replaces human verification" but "automated screening reduces the verification burden imposed on humans."

5.5 Research-integrity governance

The problem has also moved into research-integrity policy. Resnik and Hosseini (2026) argue that hallucinated citations may constitute research misconduct in certain circumstances, particularly when citations function as data in scholarly work and researchers knowingly or recklessly fail to verify them. They distinguish this from ordinary citation errors and contextual inaccuracies.

The governance implication is important: citation integrity is increasingly being treated as part of research integrity rather than as a minor formatting issue.

6. Research Challenges and Limitations

The principal challenge is the lack of direct causal evidence concerning *formatting itself*. Existing studies establish several adjacent propositions:

1. source cues can influence perceived credibility;

2. citation presence can increase trust in AI-generated content;
3. fabricated citations occur in AI outputs;
4. fabricated references can enter published literature;
5. bibliographic metadata can be manipulated; and
6. citation errors are common enough to warrant systematic verification.

These findings jointly make a formatting effect plausible, but they do not establish that changing APA-like punctuation, numeric citation style, DOI display, or bibliographic completeness independently changes trust.

A second limitation concerns ecological validity. Much existing credibility research uses online articles, websites, news stories, or AI answers rather than scholarly manuscripts. Academic readers operate within different expectations and possess specialized bibliographic knowledge.

Third, expertise is likely to moderate effects. A librarian, researcher, or graduate student may immediately recognize a malformed DOI or impossible journal-volume combination that a general reader overlooks. Lucassen and Schraagen's (2013) findings concerning topic familiarity illustrate why user knowledge should be incorporated into experimental designs.

Fourth, "trust" is multidimensional. Researchers should distinguish:

- perceived credibility;
- willingness to believe a claim;
- willingness to cite the source;
- willingness to act on the information;
- willingness to recommend the information;
- confidence in the document;
- actual verification behavior.

A person may report high credibility while still deciding to verify a citation before acting.

Finally, fabricated references are heterogeneous. Some are completely nonexistent; others combine real authors with nonexistent titles; others use genuine articles but incorrect metadata; still others cite real sources that do not support the stated claim. A single experimental category called "fake citation" therefore risks obscuring important differences.

7. Research Gaps and Future Directions

7.1 Factorial experiments isolating formatting

The most important research gap is a controlled factorial experiment that independently manipulates:

- **source validity:** genuine vs. fabricated;
- **format quality:** low vs. high conformity to disciplinary conventions;

- **bibliographic completeness:** minimal vs. detailed;
- **identifier presence:** no DOI vs. valid DOI vs. invalid DOI;
- **citation style:** author–date vs. numeric;
- **citation relevance:** relevant vs. irrelevant;
- **document quality:** professionally written vs. poorly written.

The primary dependent variable should be perceived trustworthiness, complemented by actual verification behavior.

Such a design would permit a direct estimate of whether formatting produces an interaction with fabrication. The critical prediction would be that high-format-quality fabricated citations receive higher *pre-verification* trust ratings than low-format-quality fabricated citations, but that this advantage disappears after verification.

7.2 Behavioral rather than self-report measures

Future work should measure whether participants actually click, search, or resolve citations. Ding et al. (2025) demonstrate the value of tracking citation-checking behavior rather than relying exclusively on self-reported trust.

Eye tracking could provide additional evidence. If readers inspect author names and journal titles but rarely examine DOIs or reference links, this would reveal which formatting elements function as early-stage credibility cues.

7.3 Expertise and disciplinary differences

Experiments should compare undergraduate students, graduate researchers, professional academics, librarians, and subject-matter experts. Citation conventions differ across disciplines, so an APA-formatted citation may appear more familiar to a psychologist than to a computer scientist accustomed to IEEE-style numbering.

A stronger design would therefore test whether *disciplinary congruence* matters more than absolute formatting quality. The relevant variable may not be "good formatting" but "formatting that matches the reader's learned expectations."

7.4 Cross-cultural and multilingual research

Most credibility research has been conducted in relatively narrow linguistic and cultural contexts. Citation conventions, perceptions of prestigious journals, and expectations concerning author names and institutional affiliation may vary internationally. Cross-cultural experiments could determine whether formatting-based trust effects generalize beyond English-language academic communication.

7.5 Longitudinal contamination studies

A further gap concerns the possibility of recursive contamination. Once fabricated references enter the scholarly record, later authors—and eventually AI systems trained on scholarly text—may reproduce them. Camp et al. (2025) document propagation of counterfeit citations and discuss the resulting damage to citation networks.

Future research should model citation networks as potentially contaminated systems in which false references acquire additional apparent legitimacy through repeated citation. This creates a distinction between *initial fabrication* and *secondary legitimization*. A nonexistent article cited once may look suspicious; the same nonexistent article cited by dozens of later papers may appear increasingly credible.

7.6 Designing citation systems that resist deception

Finally, scholarly interfaces should shift from *presentation of citations* toward *presentation of verification status*. A future reference interface might display whether a DOI resolves, whether metadata agree across databases, whether the cited work exists in recognized indexes, and whether the source has been retracted or corrected.

Such systems would make verification visible at the point of reading rather than requiring readers to perform it manually. They could therefore reduce the opportunity for formatting alone to substitute for evidence.

8. Conclusion

Citation formatting is conventionally treated as an editorial or stylistic concern. The evidence reviewed here suggests that this interpretation is incomplete. Citation presentation is also part of the credibility environment in which readers evaluate scholarly claims. Research on source credibility demonstrates that people use contextual and source cues; research on AI-generated answers demonstrates that citations can increase perceived trustworthiness; and recent audits show that fabricated references can imitate the bibliographic conventions of genuine scholarship with considerable sophistication.

Nevertheless, the literature does **not yet establish that citation formatting alone causally increases trust in fabricated resources**. This distinction is essential. Evidence that citations increase trust is not equivalent to evidence that APA, MLA, Vancouver, or another particular formatting convention increases trust. Similarly, evidence that fabricated citations are often correctly formatted demonstrates the *possibility of exploitation*, not the psychological magnitude of the formatting effect.

The strongest interpretation of current evidence is therefore a **credibility-cue hypothesis**: polished, complete, consistent, and discipline-congruent citation formatting may increase initial perceptions of legitimacy by reducing uncertainty, increasing processing fluency, and signaling scholarly competence. When the underlying source is fabricated, these benefits may operate

before verification. Once a reader checks the citation and discovers that the resource is nonexistent or mismatched, the credibility signal should reverse.

The practical implication is that citation formatting must never be treated as evidence of source validity. The scholarly community should distinguish presentation quality from bibliographic authenticity and claim support. Automated reference verification, DOI resolution, metadata reconciliation, and human inspection should complement—rather than be replaced by—well-formatted references. Recent evidence that fabricated references have entered peer-reviewed literature at scale makes this distinction increasingly important.

For researchers, editors, publishers, educators, and developers of AI-assisted writing systems, the central lesson is straightforward: **a citation can look scholarly without being scholarly evidence**. The next generation of research should determine precisely how much that appearance influences human judgment and under what conditions readers can be trained or technologically assisted to resist it.

References

Agarwal, A., Arafa, M., Avidor-Reiss, T., Hamoda, T. A.-A. A.-M., & Shah, R. (2023). Citation errors in scientific research and publications: Causes, consequences, and remedies. *The World Journal of Men's Health*, 41(3), 461–465. <https://doi.org/10.5534/wjmh.230001>. PMID: 37118953; PMCID: PMC10307651.

Besançon, L., Cabanac, G., Labbé, C., & Magazinov, A. (2024). Sneaked references: Fabricated reference metadata distort citation counts. *Journal of the Association for Information Science and Technology*, 75(12), 1368–1379. <https://doi.org/10.1002/asi.24896>. arXiv:2310.02192.

Camp, N. T., Bengtson, J. A., & Sandstrom, J. C. (2025). The citation catastrophe: Propagation of AI-generated counterfeit citations in scholarship. *The Journal of Academic Librarianship*, 51(4), 103065. <https://doi.org/10.1016/j.acalib.2025.103065>.

Ding, Y., Facciani, M., Joyce, E., Poudel, A., Bhattacharya, S., Veeramani, B., Aguiñaga, S., & Weninger, T. (2025). Citations and trust in LLM generated responses. *Proceedings of the AAAI Conference on Artificial Intelligence*, 39(22), 23787–23795. <https://doi.org/10.1609/aaai.v39i22.34550>. arXiv:2501.01303.

Fogg, B. J., & Tseng, H. (1999). The elements of computer credibility. In *CHI '99: Proceedings of the SIGCHI Conference on Human Factors in Computing Systems* (pp. 80–87). Association for Computing Machinery. <https://doi.org/10.1145/302979.303001>.

Go, E., Jung, E. H., & Wu, M. (2014). The effects of source cues on online news perception. *Computers in Human Behavior*, 38, 358–367. <https://doi.org/10.1016/j.chb.2014.05.044>.

Hovland, C. I., & Weiss, W. (1951). The influence of source credibility on communication effectiveness. *Public Opinion Quarterly*, 15(4), 635–650. <https://doi.org/10.1086/266350>.

Kang, H., Bae, K., Zhang, S., & Sundar, S. S. (2011). Source cues in online news: Is the proximate source more powerful than distal sources? *Journalism & Mass Communication Quarterly*, 88(4), 719–736. <https://doi.org/10.1177/107769901108800403>.

Klein, J., & Wager, E. (2019). Formatting references for scientific manuscripts. *Journal of Medical Library Association*, 107(2), 283–286. [Bibliographic discussion of Vancouver and Harvard reference systems.]

Lucassen, T., & Schraagen, J. M. (2013). The influence of source cues and topic familiarity on credibility evaluation. *Computers in Human Behavior*, 29(4), 1387–1392. <https://doi.org/10.1016/j.chb.2013.01.036>.

Resnik, D. B., & Hosseini, M. (2026). Hallucinated citations produced by generative artificial intelligence may constitute research misconduct when citations function as data in scholarly papers. *Accountability in Research*, 1–8 [online ahead of print]. <https://doi.org/10.1080/08989621.2026.2645390>. PMID: 41833014; PMCID: PMC13051339.

Topaz, M., Roguin, N., Gupta, P., Zhang, Z., & Peltonen, L.-M. (2026). Fabricated citations: An audit across 2·5 million biomedical papers. *The Lancet*, 407(10541), 1779–1781. [https://doi.org/10.1016/S0140-6736\(26\)00603-3](https://doi.org/10.1016/S0140-6736(26)00603-3). PMID: 42107362.

Walters, W. H., & Wilder, E. I. (2023). Fabrication and errors in the bibliographic citations generated by ChatGPT. *Scientific Reports*, 13, 14045. <https://doi.org/10.1038/s41598-023-41032-5>. PMID: 37679503; PMCID: PMC10484980.

Note on evidentiary scope

The paper deliberately distinguishes **direct evidence** from **theoretical inference**. The strongest direct evidence currently concerns (a) the effects of source cues on credibility, (b) the effects of citation presence and validity on trust in AI-generated responses, and (c) the prevalence and plausibility of fabricated references. Direct experiments isolating *citation formatting itself* as an independent causal variable remain comparatively scarce. Accordingly, claims about formatting increasing trust in fabricated resources are presented as a theoretically grounded hypothesis and research agenda rather than as an established empirical fact. This distinction is necessary to avoid reproducing the very evidentiary problem examined in the paper.